Neuro-symbolic verifier for deterministic AI outputs.

Smart as an LLM. Precise as math. Cost effective.

Françoise

The LLM compliance trap

LLMs are probabilistic by design. Even top reasoning models still hallucinate.
In enterprise compliance, one wrong parameter is enough: broken UI, legal exposure, fail QA.

The default fix is expensive and unstable:
more prompts, more retries, more token burn.
Each repair loop is a new billed model call.

You can't solve a math problem with more probability

Source: Glean / OpenAI benchmark summary;
OpenAI API Pricing; Self-Refine / prompt chaining literature

Hybrid intelligence

Proprietary Neuro-Symbolic Engine

Contextual Translation:
Our AI maps abstract, complex rules into programmable constraints.
Deterministic Verification Core:
An algorithmic layer mathematically proves 100% compliance with these constraints.
Surgical Resolution: Only logically conflicting edge cases are routed back to the LLM.

We use LLMs for meaning, and algorithms for truth

Precision at a fraction of the cost

	Standard LLM workflow Generate → retry → patch → retry	Françoise Flow Generate ↔ snap → production ready
Token cost	Every repair loop is a new billed model call	**Structural fixes use zero LLM tokens**
Fixing cost	40–60% dev time wasted on manual QA	**Zero manual QA for structural rules**
Reliability	Probabilistic, with hallucination risk	**Deterministic on structural rules**
Failure mode	Errors cascade across iterations	**One precise fallback prompt**

OpenAI API Pricing; Glean on enterprise hallucinations; Self-Refine / prompt chaining

UI compliance first

UI is where creative intent collides with hard rules: WCAG, spacing, contrast, tokens, components. **That makes it the perfect wedge for a neuro-symbolic verifier.**

The pain is already budgeted:

The Risk: EAA compliance is legally enforced since June 2025 (fines up to €600k in EU). ADA fines reach $150k per violation in the US.
The ROI: At $45k/year, a single avoided fine pays for 3 to 13 years of Françoise.
The Hidden Cost: Manual QA is slow and error-prone. Replacing just a fraction of manual UI testing covers our annual license instantly.

Françoise turns compliance from a QA bottleneck into an automatic verification layer

EAA country penalties; ADA fines

The competitive landscape

Existing approach	What it does	Why it breaks
Manual QA	Humans check outputs	Slow, expensive, doesn't scale
Linters / checkers	Catch simple rule violations	Don't fix the output
LLMs alone	Generate fast	Break design systems, no guarantees
Françoise	**Understands intent, then enforces rules deterministically** **Built for enterprise-grade certainty**	

Our edge:
Deterministic Accuracy: AI handles intent, algorithms enforce 100% rule compliance.
Zero Manual QA: We bridge AI directly to enterprise design systems.
Deep Enterprise Moat: Deep CI/CD integration and custom Design System mapping create massive switching costs.

Early market pull

Converting Pilots to ARR: $5K Paid Pilots → $45K Annual Contracts.

Paid pilot (5k$):
Tier-1 global e-commerce platform

Active POCs:
inDrive, SalaryHub, AffArts

Advanced pipeline:
Microsoft, TripleTen, Miro, Semrush

25+ enterprise / scale-up leads in 12 weeks

They are testing a way to get deterministic outputs where LLMs alone break down.

Business Model

Target: Enterprises with 50-200 product designers in regulated industries.

Land: Paid Pilot ($5K)

8-week proof of concept on client's live data.
Bypasses enterprise procurement delays.

Expand: Annual Platform License

$45K / year / 50 seats

+ $10K for CI/CD Pipeline Integration
+ $20K for On-Premise / Private Cloud Deployment

Path to $1B+ Valuation:

1,000 enterprise clients × $65K ACV (License + On-prem) = $65M ARR.
(Capturing just 2% of the 50k+ global regulated enterprises).

85% gross margin

achieved by offloading structural verification from billed tokens to local deterministic code

From UI wedge to verification infrastructure

We start with UI and accessibility compliance because the pain is acute, the rules are strict, and the ROI is immediate.

Then the same architecture expands into:

Front-end code verification
Agent output validation
High-stakes workflows in regulated industries

The underlying market is much larger than design tooling: AI inference is a $106B market in 2025 and growing fast.

Source: MarketsandMarkets

Raising $500k pre-seed

to $250K ARR that triggers our $1.5M Seed round
(in active talks with Tier-1 EU & US funds)

Team

Daria Kamyshina, CEO & Enterprise Sales.
ex-Head of Product (EdTech). Founder of AI agency,
closed enterprise deals with tech giants like Yandex and Avito.

Victor Pisarevskii, CTO & AI Architecture
MSc Applied Math (Honors). Architected high-load FinTech
solutions scaling to 10M+ active users.

daria@francoise.design | meet.francoise.design | francoise.design