Neuro-symbolic verifier for deterministic AI outputs.

Smart as an LLM. Precise as math. Cost effective.

Françoise

The LLM compliance trap

LLMs are probabilistic by design. Even top reasoning models still hallucinate.
In enterprise compliance, one wrong parameter is enough: broken UI, legal exposure, fail QA.

The default fix is expensive and unstable:
more prompts, more retries, more token burn.
Each repair loop is a new billed model call.

You can't solve a math problem with more probability

Source: Glean / OpenAI benchmark summary;
OpenAI API Pricing; Self-Refine / prompt chaining literature

Hybrid intelligence:

semantics from AI, certainty from code

Françoise sits between the LLM and the user.

AI

Semantic mapping: a low-cost model understands intent
Algorithmic snapping: deterministic code enforces the rules exactly
Smart prompting: only unresolved logic goes back in one targeted prompt

We use LLMs for meaning, and algorithms for truth

Precision at a fraction of the cost

	Standard LLM workflow Generate → retry → patch → retry	Françoise Flow Generate ↔ snap → production ready
Token cost	Every repair loop is a new billed model call	**Structural fixes use zero LLM tokens**
Fixing cost	40–60% dev time wasted on manual QA	**Zero manual QA for structural rules**
Reliability	Probabilistic, with hallucination risk	**Deterministic on structural rules**
Failure mode	Errors cascade across iterations	**One precise fallback prompt**

OpenAI API Pricing; Glean on enterprise hallucinations; Self-Refine / prompt chaining

UI compliance first

UI is where creative intent collides with hard rules: WCAG, spacing, contrast, tokens, components. **That makes it the perfect wedge for a neuro-symbolic verifier.**

The pain is already budgeted:

EAA: enforcement began in June 2025; penalties can reach €600k in Spain, €250k in France, €100k in Germany
ADA: fines can reach $75k for a first violation and $150k for repeat violations
Manual QA: slow, expensive, and still error-prone

Françoise turns compliance from a QA bottleneck into an automatic verification layer

EAA country penalties; ADA fines

The competitive landscape

Existing approach	What it does	Why it breaks
Manual QA	Humans check outputs	Slow, expensive, doesn’t scale
Linters / checkers	Catch simple rule violations	Don’t fix the output
LLMs alone	Generate fast	Break design systems, no guarantees
Françoise	**Understands intent, then enforces rules deterministically** **Built for enterprise-grade certainty**	

Our edge:
We remove manual QA from structural verification
We don’t use LLMs as a source of truth for math
We bridge AI-generated output to real enterprise design systems

Early market pull

High-stakes teams are already testing this approach in real workflows

Paid pilot:
Tier-1 global e-commerce platform

Active POCs:
inDrive, SalaryHub, AffArts

Advanced pipeline:
Microsoft, Wolt, Miro, Semrush, Emirates NBD

25+ enterprise / scale-up leads in 12 weeks

They are testing a way to get deterministic outputs where LLMs alone break down.

Business Model

Target: Enterprises with 50-200 product designers in regulated industries.

Land: Paid Pilot ($5K)

8-week proof of concept on client's live data.
Bypasses enterprise procurement delays.

Expand: Annual Platform License

$45K / year / 50 seats

+ $10K for CI/CD Pipeline Integration
+ $20K for On-Premise / Private Cloud Deployment

85% gross margin

From UI wedge to verification infrastructure

We start with UI and accessibility compliance because the pain is acute, the rules are strict, and the ROI is immediate.

Then the same architecture expands into:

Front-end code verification
Agent output validation
High-stakes workflows in regulated industries

The underlying market is much larger than design tooling: AI inference is a $106B market in 2025 and growing fast.

Source: MarketsandMarkets

Raising $500k pre-seed

to $250K ARR that triggers our $1.5M Seed round
(in active talks with Target Global & Enable Ventures)

Team

Daria Kamyshina, CEO
ex-Head of Product, ex-founder AI implementation agency

Victor Pisarevskii, CTO
ex-Tech Lead at Tier-1 FinTech; MSc in Applied Math with Honors

daria@francoise.design | meet.francoise.design | francoise.design